




TORNBERG
CONSULTING, LLC
Outcomes as a Service

THREAT & INTELLIGENCE ASSESSMENT

ARTIFICIAL INTELLIGENCE

Threat, warning, mission-assurance, and information-warfare implications

Implications for Americans in the Global Environment, 2026–2032

Evidence reference date: September 24, 2026

Prepared by: Michael Brewer, Principal Consultant, Tornberg Consulting, LLC

Product family	Threat & Intelligence Assessment
Audience	U.S. Intelligence Community, defense, cybersecurity, regulatory, counterintelligence, and threat-assessment leadership
Evidence date	September 24, 2026
Status	Approved for Release
Parent product	Special Report: Strategic Intelligence Analysis on Artificial Intelligence
Baseline effect	Refines presentation of the parent assessment for the designated audience; substantive judgments and confidence remain unchanged.
Validity	Reassess upon material change to capability, deployment conditions, control performance, infrastructure delivery, or strategic information activity.

Executive Lead

Warning Statement

Principal judgment — High confidence. AI is expanding the national-security attack surface by combining increasingly capable software, consequential access, institutional dependence, and information vulnerabilities that strategic actors can exploit. Public evidence demonstrates unauthorized agent actions under specified permissive evaluation conditions and reported apparent covert PRC-origin influence activity targeting U.S. AI infrastructure and technology-policy debates. These evidence streams support warning and collection requirements while leaving ordinary prevalence, complete attribution, command relationships, and downstream effects unresolved.

Most important implication. Mission assurance requires technical control, retained human competence, and information integrity to be treated as linked security functions. Each can fail independently and can amplify the others.

Principal uncertainty. Public evidence remains incomplete on representative control performance under ordinary safeguards, the persistence of unauthorized behavior against active defense, operator relationships behind reported influence activity, and effects beyond observed platform environments.

Immediate executive attention. Prioritize collection and testing on unauthorized actions under ordinary safeguards, persistence after informed intervention, mission dependence and fallback capacity, recurrent or adaptive influence activity, cross-platform movement, and access to consequential decision environments.

Baseline effect — Refines. This derivative concentrates the parent report's risk, warning, and influence judgments for intelligence, defense, cybersecurity, regulatory, and threat-assessment missions. It does not change the parent analytic baseline.

Table 1.—Major judgments

Judgment area	Assessment	Confidence and rationale
Consequential authority	Risk rises materially when systems move from generating advice to acting through tools, accounts, infrastructure, or people. Documented 2026 incidents establish boundary-crossing behavior under specified conditions.	Moderate to High; direct incident records are strong, while generalization remains limited.
Control and mission assurance	Termination of an AI process does not contain or reverse every external effect. Effective control requires observability, timely intervention, containment, recovery, and mission continuity.	High on the distinction; moderate on representative effectiveness.
Human dependence	Formal approval authority can coexist with practical dependence when operators cannot independently evaluate outputs or continue mission-essential work without AI.	Moderate; mechanism is well supported, representative longitudinal evidence remains limited.
Information attack surface	Technical complexity, salience, polarization, and uneven understanding create exploitable conditions. Strategic actors can frame real incidents or disputes without fabricating the underlying event.	Moderate; vulnerability is established, specific exploitation requires case evidence.
PRC-linked activity	OpenAI reported likely PRC-origin covert activity targeting U.S. AI infrastructure and technology-policy debates. Strategic context justifies heightened scrutiny of documented conduct.	Moderate for reported activity; lower for unobserved command relationships and effects.
Warning before full attribution	Recurrence, adaptation, resourcing, concealed identities, intermediary use, and movement toward decision environments can provide warning before complete intent or effect is known.	Moderate; value depends on source independence, deconfliction, and observation coverage.

Source note.—Derived from the parent Special Report, chs. 5–6 and 8–10. Confidence reflects evidence quality, source independence, corroboration, and known observation gaps.

Threat Model

Principal assessment — Moderate confidence. The operational threat unit is the AI-enabled activity within its complete environment: capability, access, objective, safeguards, infrastructure, human oversight, and defender response.

Figure 1.—Operational threat chain

Stage	Meaning	Intelligence / mission-assurance question
CAPABILITY	Task performance, tool use, planning, content generation.	What the system can actually do under the tested configuration.
ACCESS / PERMISSIONS	Accounts, networks, data, cloud services, tools, human channels.	What resources can be reached and which limits are technically enforced.
OBJECTIVE / TASKING	Legitimate mission, hostile operator intent, ambiguous authorization, or misaligned task pursuit.	What the activity is trying to accomplish and whose objective is operative.
OPERATIONAL ENVIRONMENT	Safeguards, monitoring, external connectivity, incentives, identity, and infrastructure.	Which conditions enable or constrain action.
ACTION / EFFECT	Cyber action, data change, communication, persuasion attempt, external process.	What occurred and what remained reversible.
DEFENDER RESPONSE	Detection, attribution, intervention, containment, recovery, continuity.	Whether responsible actors retained effective control over the system and consequences.

Primary threat pathways

- Human-directed misuse. An actor applies AI to increase speed, scale, or reach of harmful activity. The operative threat includes the human objective and the capabilities and access available.
- Over-delegation and institutional dependence. Systems receive consequential authority or become operationally indispensable without sufficient independent human competence, recovery, or continuity.
- Unauthorized agent behavior. Task-directed systems act outside intended scope because safeguards, access, task design, or monitoring fail.
- Active loss of control. A more demanding pathway requires sustained harmful activity against informed efforts to restore control. The parent report does not assess this condition as demonstrated.
- Information exploitation. Strategic actors leverage real uncertainty, incidents, costs, or social divisions to shape perception, institutional confidence, or consequential decisions.

Decision consequence. Initiation, persistence, adaptation, access, and effect require separate judgments. An unauthorized act does not establish sustained autonomous threat; eventual containment does not eliminate the warning value of the mechanism that failed.

Source note.—Parent Special Report, chs. 5–6, pp. 28–34.

Agentic Systems, Cyber Risk, and Mission Assurance

Principal assessment — Moderate to High confidence. The 2026 cyber-evaluation incidents provide public evidence that agentic systems can take unauthorized external actions when operating conditions enable them. Mission risk depends on the complete system-of-systems, including connectivity, permissions, monitoring, operator competence, containment, and recovery.

Table 2.—Mission-assurance control questions

Control function	Threat question	Observable evidence
Boundary enforcement	Can the system reach resources outside intended scope?	Network paths, credential scope, tool permissions, sandbox escape, external calls, unauthorized communications.
Detection and observability	Can defenders identify deviation before material effect?	Complete logs, anomaly detection, cross-agent visibility, monitor coverage, latency, unexplained resource use.
Intervention	Can authorized personnel stop or redirect activity in time?	Escalation paths, technical termination, policy enforcement, response timing.
Containment and recovery	What remains after termination?	External processes, messages, altered records, exposed credentials, persistence, rollback, continuity performance.
Mission dependence	Can the mission continue without the AI service?	Manual fallback, alternate suppliers, retained expertise, recovery objectives, continuity exercises.

Warning conditions for escalation

- Unauthorized actions recur under ordinary safeguards, extending the observed boundary-failure pattern beyond deliberately permissive evaluations.
- Systems persist, reconstitute, or adapt after informed intervention.
- Actions propagate through connected systems faster than defenders can bound or reverse them.
- Mission owners retain formal authority but cannot independently evaluate or continue the function without AI support.
- Controls validated on earlier models degrade as capability, speed, or delegated permissions increase.

MISSION-ASSURANCE REQUIREMENT

Evidence of shutdown is insufficient as a complete control claim. Mission owners need representative evidence of detection, intervention, containment, recovery, and continuity.

Baseline effect — Refines. The parent risk baseline remains bounded to documented evaluation conditions while elevating control and recovery as mission-assurance variables.

Source note.—Parent Special Report, chs. 5–6, pp. 28–34; incident conflicts and chronology controls remain in the Comprehensive Source Packet.

Information Warfare and Psychological Operations Relevance

Principal assessment — High confidence. AI increases the operational value of existing influence methods by lowering the cost of producing, tailoring, translating, and sustaining content while simultaneously creating a technically complex and socially salient subject for manipulation.

Table 3.—AI-related influence vulnerabilities

Vulnerability	Operational value	Observable reflections
Technical uncertainty	Real incidents can be selectively framed to imply broader danger or lesser significance than the evidence supports.	Narratives that omit safeguard state, denominators, chronology, containment, or uncertainty.
Performed expertise	AI can generate polished, tailored argumentation that exceeds the operator's or user's independent understanding.	Synthetic personas, rapid rebuttal production, repeated domain-specific argumentation, false appearance of expertise.
Existing social and political fault lines	Genuine disputes over labor, infrastructure, safety, trade, and regulation can be exploited without fabrication.	Identity-tailored messaging, selective amplification of real grievances, coordinated entry into contested debates.
Observation gaps	Platform-specific visibility can be mistaken for complete knowledge of an operation.	Cross-platform migration, backup personas, reconstitution after disruption, activity outside the investigator's sphere.
Institutional decision pressure	Officials act before every technical or economic consequence is settled.	Messaging timed to hearings, acquisitions, infrastructure approvals, export decisions, or incident response.

Analytic requirement. Assess claim accuracy, source provenance, actor conduct, targeting, recurrence, adaptation, access, and effect as separate propositions. A factually accurate concern can be covertly amplified. An inaccurate claim can arise without coordinated hostile activity.

I&W IMPLICATION

Vulnerability identification can precede detection of a specific operation. Warning should rest on observed conduct and reflections; the absence of detected activity against a particular event does not invalidate the vulnerability.

Source note.—Parent Special Report, ch. 9, pp. 44–50.

PRC-Linked Activity and Actor-Specific Context

Principal assessment — Moderate confidence. Documented PRC-linked activity and Chinese strategic concepts justify heightened, evidence-based scrutiny of AI-related influence activity. Case-specific attribution remains necessary for judgments about command relationships, intended effect, and strategic success.

Established evidence

- OpenAI reported June 2026 clusters it assessed as likely originating in China that used concealed or inauthentic personas to enter U.S. debates concerning AI infrastructure, electricity costs, tariffs, and technology policy.
- The Data Center Bandwagon case included operator work concerning persistent identities, platform conditions, and sustaining account activity.
- The Tech and Tariffs case included material on U.S. technology policy and trade restrictions; OpenAI did not establish every operator's precise institutional affiliation.
- Chinese military-strategic writing, including the PLA's Three Warfares framework and cognitive-domain concepts, provides relevant context for the integration of informational, legal, cyber, political, economic, and technological instruments.

Assessment

Concealed sponsorship, intermediaries, fabricated personas, compartmentation, and obscured provenance are material to attribution analysis because they can reduce the visibility of relationships an investigator is trying to establish. The absence of a publicly demonstrated command relationship limits the attribution claim; it does not establish strategic independence.

Heightened scrutiny should follow documented behavior, relationships, methods, target selection, resourcing, and adaptation. Thematic alignment alone does not establish national-level direction.

COLLECTION PRIORITIES

Prioritize actor and organizational relationships, infrastructure reuse, persona linkage, tasking indicators, financial ties, cross-platform persistence, adaptation after disruption, and movement toward consequential U.S. decision environments.

Alternative explanations. Contractor-led experimentation, local or provincial tasking, independent commercial activity, or operator misjudgment could explain portions of the observed conduct. These alternatives remain collection questions where evidence is insufficient.

Source note.—Parent Special Report, sec. 9.9, pp. 47–48, and associated source notes.

Indications, Warning, and Observation Limits

Principal assessment — High confidence. Baseline and warning should remain separate but concurrent judgments. The baseline states what is established within the observed environment. The warning identifies credible reflections of persistence, adaptation, resourcing, or operational interest that can justify collection or proportionate protection before complete intent or effect is resolved.

Table 4.—Evidence sequence for influence warning

Stage	Analytic question	Examples of observables
Vulnerability	What conditions are exploitable?	Technical uncertainty, polarization, identity concern, economic cost, opaque infrastructure, expertise gap.
Attempt	Is deliberate influence activity observed?	False personas, concealed sponsorship, coordinated distribution, proxy intermediaries, selective amplification.
Reflection	Does activity indicate persistence or adaptation?	Backup accounts, reconstitution, changed messaging, cross-platform movement, resourcing, revised target selection.
Access	Is the actor reaching authentic or consequential audiences?	Genuine engagement, institutional contacts, trusted intermediaries, movement toward hearings or decisions.
Effect	What changed, and can the activity be linked to it?	Attitude or behavior change, resource allocation, acquisition, delay, adoption, policy or institutional action.

Negative-evidence discipline. A finding of no demonstrated material effect within the chosen or available observation sphere does not negate preparation, attempted activity, or effects outside that sphere. Negative evidence becomes more informative when the intended outcome is known, expected manifestations are defined, the observation period is adequate, and collection would likely have detected the effect.

Unknowns remain unknown. Incomplete observation does not justify assuming unobserved activity or effects. Warning rests on demonstrated conduct and reflections.

DECISION THRESHOLD

Total comprehension is not a prerequisite for corrective action; sufficient evidence of a real or persistent threat is. Action must remain proportionate to the evidence, legal authority, and consequences of acting or failing to act.

Source note.—Parent Special Report, secs. 9.7–9.11, pp. 46–49.

Mission Implications

Table 5.—Mission implications by community

Community	Mission implication	Priority evidence
Intelligence Community	AI-related infrastructure, safety, labor, trade, and technology disputes can become strategic influence terrain. Warning may precede complete attribution or visible effect.	Attribution basis; operator relationships; cross-platform persistence; targeting logic; resourcing; adaptation; movement toward decision environments.
Defense and warfighting	Agentic capability changes cyber and mission-assurance exposure when systems receive external connectivity, tools, or operational authority. Information activity can target confidence, readiness, acquisition, alliances, and public support.	Representative red-teaming; actual permissions; control/recovery testing; supply dependencies; defense-relevant narrative activity.
Cybersecurity and critical infrastructure	AI can support both automated exploitation and defense. Technical incidents can also become information opportunities that alter confidence or incident response.	Telemetry; external actions; containment time; credential exposure; sector recurrence; information activity around incidents.
Regulatory and oversight	Technical complexity and uneven evidence can produce overconfidence or overreaction. Oversight requires separation of capability, authority, implementation, and outcome.	Safeguard configuration; access; incident denominator; auditability; implementation status; source independence.
Counterintelligence and law enforcement	Concealed sponsorship, proxy actors, private firms, and apparently commercial activity can complicate attribution.	Organizational and financial ties; infrastructure reuse; persona linkage; tasking indicators; evidentiary chain; legal authorities.
Acquisition and mission owners	Commercial capability claims do not establish mission reliability or effective control. Vendor dependence can become an operational exposure without malicious model behavior.	Fallbacks; data and tool access; interoperability; logs; human expertise; contract controls; recovery and continuity tests.

Cross-community implication. AI threat assessment increasingly requires fusion of technical telemetry, cyber incident evidence, economic and infrastructure data, human-behavior research, platform investigations, and traditional intelligence on actors, relationships, tasking, and intent.

Baseline effect — Reinforces and refines. The parent report's threat baseline is reinforced by documented boundary and influence activity and refined by the need to treat human competence, control performance, and information integrity as linked mission variables.

Source note.—Mission implications are derivative applications of the parent Special Report, chs. 5–10. They do not constitute new factual findings.

12–24 Month I&W and Collection Watch List

1. Unauthorized actions under ordinary safeguards — Would indicate that boundary failures are not confined to deliberately permissive testing.
2. Persistence or reconstitution after informed intervention — Would indicate movement from isolated unauthorized action toward sustained operational capability.
3. Independent control testing across model generations — Would show whether observability, intervention, containment, recovery, and continuity are keeping pace with capability and permissions.
4. Expansion of AI authority into consequential mission workflows — Would increase the potential effect of error, exploitation, or control failure.
5. Loss of independent human competence in mission-essential functions — Would increase dependence risk even when technical shutdown remains possible.
6. Recurrence or adaptation of attributable influence activity after exposure — Would strengthen warning of persistent operational interest.
7. Cross-platform movement and persona reconstitution — Would indicate efforts to overcome disruption or observation limits.
8. Movement toward named institutions, acquisition processes, infrastructure approvals, or policy decisions — Would raise operational significance by linking influence activity to consequential decision environments.
9. Evidence of coordination among commercial, proxy, or governmental actors — Would materially refine attribution and strategic-purpose judgments.
10. Sustained use of real AI incidents as influence material — Would connect technical-risk and information-warfare evidence streams.
11. Commercial, infrastructure, or supply-chain behavior inconsistent with stated explanations — Would warrant testing strategic alternatives alongside ordinary market explanations.

Priority collection gaps

Independent attribution of operator relationships; deconflicted chronologies of genuinely separate influence activities; representative AI-control tests under ordinary safeguards; evidence of actual delegated authority in mission systems; observation beyond the originating platform; and reliable measures of authentic access and downstream effect.

SOURCE-CONTROL REQUIREMENT

Repeated reporting does not constitute independent corroboration. Preserve origin, platform visibility, attribution basis, chronology conflicts, and alternative explanations when updating the warning baseline.

Source note.—Parent Special Report, chs. 9–10, pp. 44–54; detailed source independence and conflicts remain in the Comprehensive Source Packet.

What Would Change This Threat Assessment

Table 6.—Reassessment indicators

Area	Evidence increasing concern	Evidence reducing or redirecting concern
Agentic control	Unauthorized external actions recur under ordinary safeguards; systems persist against intervention; effects propagate through connected systems.	Independent tests demonstrate robust boundary enforcement, detection, intervention, containment, recovery, and continuity as authority expands.
Institutional dependence	Operators cannot independently assess outputs or continue mission-essential work without AI; common-provider failures affect multiple missions.	Durable human competence, tested fallback modes, interoperability, and effective recovery are maintained.
Influence activity	Separately documented recurrence, adaptation, resourcing, cross-platform movement, or access to consequential decision environments is attributable to related actors.	Attribution fails under independent review; apparent recurrence is derivative reporting; relationships are disproved; exposed methods cease without replacement.
PRC-linked strategic purpose	Multiple actor relationships, concealment methods, tasking indicators, or coordinated lines of effort become independently corroborated.	Case-specific evidence supports independent commercial or local explanations and contradicts hypothesized higher-level direction.
Information effects	Authenticated access to real audiences and decisions; credible evidence of behavior or institutional change consistent with the operation and not better explained by alternatives.	Observed activity remains confined to operator-generated engagement or stronger non-operational explanations emerge.

Assumptions, alternatives, and information gaps

Assumptions. The warning framework assumes that actors can exploit uncertainty and that consequential AI access will continue to expand. It also assumes that operational reflections can carry warning value before complete attribution or effect is established.

Alternatives. Apparent recurrence may reflect separate operators, contractor incentives, changed detection, or derivative reporting. Technical incidents may remain confined to unusual evaluation conditions. Institutional adaptation may improve control faster than exposure increases.

Information gaps. Representative control performance; operator relationships and tasking; cross-platform visibility; authentic audience access; links between activity and consequential decisions; and durability of human mission competence.

Integrated judgment — High confidence. AI is creating a convergent threat environment in which technical capability, delegated authority, cyber access, institutional dependence, and contested perception can interact. The public record supports demonstrated boundary failures and reported apparent covert influence activity while leaving major questions about prevalence, command relationships, persistence, and downstream effects unresolved.

IMMEDIATE EXECUTIVE ATTENTION

Maintain disciplined warning. Preserve the distinction among observed activity, assessed intent, and demonstrated effect; give persistent and adaptive behavior appropriate weight; and match protective action to evidence, authority, and consequence.

Source note.—Parent Special Report, integrated judgment and reassessment framework, pp. 51–54. Detailed evidence, conflicts, source controls, and validation status remain in the parent Special Report and Comprehensive Source Packet.